

FoML Hackathon 2023

Name-Sanyam Kaul, Roll No-CS23MTECH14011

Name-Mayuresh Rajesh Dindorkar, Roll No-CS23MTECH14007

```
In [57]: import pandas as pd
import numpy as np
from sklearn.metrics import f1_score, accuracy_score
from sklearn.impute import KNNImputer
import matplotlib.pyplot as plt
from random import randint
import seaborn as sns
from sklearn.model_selection import RandomizedSearchCV
from sklearn.ensemble import GradientBoostingClassifier

import warnings
warnings.filterwarnings("ignore")
```

Reading the dataset

```
In [58]: train_set = pd.read_csv('iith_foml_2023_train.csv')
test_set = pd.read_csv('iith_foml_2023_test.csv')
```

```
In [59]: print('TrainSet: ', train_set.shape)
print('TestSet: ', test_set.shape)
```

```
TrainSet: (994, 25)
TestSet: (426, 24)
```

```
In [60]: train_set.head()
```

```
Out[60]:
```

	Feature 1 (Discrete)	Feature 2 (Discrete)	Feature 3 (Discrete)	Feature 4 (Discrete)	Feature 5 (Discrete)	Feature 6 (Discrete)	Feature 7 (Discrete)	Feature 8 (Discrete)	Feature 9
0	1404	12	64	14	3	1	1	1	110.502
1	909	0	235	32	1	1	1	1	-40.448
2	654	3	175	2	1	1	1	1	-27.445
3	1372	12	382	14	2	0	1	0	0.001
4	786	3	199	2	1	0	1	0	0.001

5 rows × 25 columns

```
In [61]: test_set.head()
```

Out[61]:

	Feature 1 (Discrete)	Feature 2 (Discrete)	Feature 3 (Discrete)	Feature 4 (Discrete)	Feature 5 (Discrete)	Feature 6 (Discrete)	Feature 7 (Discrete)	Feature 8 (Discrete)	Feature 9
0	146	12	42	14	7	1	1	1	118.004
1	35	0	12	5	0	0	1	0	0.001
2	1018	8	259	2	1	1	1	1	NaN
3	383	7	117	5	1	1	1	1	53.002
4	1216	7	40	5	2	0	1	4	0.005

5 rows × 24 columns

```
In [62]: X_train = train_set.drop('Target Variable (Discrete)', axis=1)
Y_train = train_set['Target Variable (Discrete)']
X_test = test_set
print('X_train: ', X_train.shape)
print('Y_train: ', Y_train.shape)
```

```
X_train: (994, 24)
Y_train: (994,)
```

```
In [63]: def pred_and_save_to_csv(X, clf, filename):
# Make predictions on the test set
y_pred = clf.predict(X)
# Create a DataFrame with "sequence_no" and "prediction" columns
results_df = pd.DataFrame({'id': X.index + 1, 'Category': y_pred})

# Save the DataFrame to a CSV file
results_df.to_csv(filename, index=False)
```

Analyzing Training Set: -

```
In [64]: X_train.isnull().sum()
```

```
Out[64]: Feature 1 (Discrete)      0
        Feature 2 (Discrete)    0
        Feature 3 (Discrete)    0
        Feature 4 (Discrete)    0
        Feature 5 (Discrete)    0
        Feature 6 (Discrete)    0
        Feature 7 (Discrete)    0
        Feature 8 (Discrete)    0
        Feature 9                14
        Feature 10               1
        Feature 11               1
        Feature 12               1
        Feature 13               1
        Feature 14               1
        Feature 15               72
        Feature 16              669
        Feature 17              546
        Feature 18              330
        Feature 19 (Discrete)    0
        Feature 20 (Discrete)    0
        Feature 21 (Discrete)    0
        Feature 22 (Discrete)    0
        Feature 23 (Discrete)    0
        Feature 24               1
        dtype: int64
```

```
In [65]: X_test.isnull().sum()
```

```
Out[65]: Feature 1 (Discrete)      0
        Feature 2 (Discrete)    0
        Feature 3 (Discrete)    0
        Feature 4 (Discrete)    0
        Feature 5 (Discrete)    0
        Feature 6 (Discrete)    0
        Feature 7 (Discrete)    0
        Feature 8 (Discrete)    0
        Feature 9                4
        Feature 10               0
        Feature 11               0
        Feature 12               0
        Feature 13               0
        Feature 14               2
        Feature 15              31
        Feature 16             279
        Feature 17             225
        Feature 18             114
        Feature 19 (Discrete)    0
        Feature 20 (Discrete)    0
        Feature 21 (Discrete)    0
        Feature 22 (Discrete)    0
        Feature 23 (Discrete)    0
        Feature 24               0
        dtype: int64
```

We can observe that more than half of the the data points in the feaure columns 'Feature 16' and 'Feature 17' are missing.

Hence, dropping the columns

```
In [66]: X_train.drop(['Feature 16', 'Feature 17'], axis=1, inplace=True)
        X_test.drop(['Feature 16', 'Feature 17'], axis=1, inplace=True)
```

Imputing the missing values with KnnImputer

```
In [67]: knn_imputer = KNNImputer(n_neighbors=5)
X_train = pd.DataFrame(knn_imputer.fit_transform(X_train), columns=X_train.columns)
X_test = pd.DataFrame(knn_imputer.transform(X_test), columns=X_test.columns)
```

```
In [68]: X_train.isnull().sum()
```

```
Out[68]: Feature 1 (Discrete)    0
Feature 2 (Discrete)    0
Feature 3 (Discrete)    0
Feature 4 (Discrete)    0
Feature 5 (Discrete)    0
Feature 6 (Discrete)    0
Feature 7 (Discrete)    0
Feature 8 (Discrete)    0
Feature 9                0
Feature 10               0
Feature 11               0
Feature 12               0
Feature 13               0
Feature 14               0
Feature 15               0
Feature 18               0
Feature 19 (Discrete)    0
Feature 20 (Discrete)    0
Feature 21 (Discrete)    0
Feature 22 (Discrete)    0
Feature 23 (Discrete)    0
Feature 24               0
dtype: int64
```

```
In [69]: X_test.isnull().sum()
```

```
Out[69]: Feature 1 (Discrete)    0
Feature 2 (Discrete)    0
Feature 3 (Discrete)    0
Feature 4 (Discrete)    0
Feature 5 (Discrete)    0
Feature 6 (Discrete)    0
Feature 7 (Discrete)    0
Feature 8 (Discrete)    0
Feature 9                0
Feature 10               0
Feature 11               0
Feature 12               0
Feature 13               0
Feature 14               0
Feature 15               0
Feature 18               0
Feature 19 (Discrete)    0
Feature 20 (Discrete)    0
Feature 21 (Discrete)    0
Feature 22 (Discrete)    0
Feature 23 (Discrete)    0
Feature 24               0
dtype: int64
```

Calculating correlation among feature columns

```
In [70]: # With the following function we can select highly correlated features
# it will remove the first feature that is correlated with anything other feature

def correlation(dataset, threshold):
    col_corr = set() # Set of all the names of correlated columns
```

```

corr_matrix = dataset.corr()
for i in range(len(corr_matrix.columns)):
    for j in range(i):
        if abs(corr_matrix.iloc[i, j]) > threshold: # we are interested in absolute values
            colname = corr_matrix.columns[i] # getting the name of column
            col_corr.add(colname)
return col_corr


```

```

In [71]: # Generate a correlation matrix
correlation_matrix = X_train.corr()

# Create a heatmap to visualize the correlation matrix
plt.figure(figsize=(12, 10))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f", linewidths=1)
plt.title('Correlation Matrix')
plt.show()

```



Dropping the columns having correlation > 0.70

```

In [72]: corr_features = correlation(X_train, 0.70)
corr_features

```

```

Out[72]: {'Feature 11',
          'Feature 13',
          'Feature 14',
          'Feature 23 (Discrete)',
          'Feature 3 (Discrete)'}

```

```
In [73]: X_train.drop(corr_features,axis=1, inplace=True)
X_test.drop(corr_features,axis=1, inplace = True)
```

```
In [74]: print(X_train.shape)
print(X_test.shape)
```

```
(994, 17)
```

```
(426, 17)
```

Gradient boost

```
In [75]: param_dist = {
    'n_estimators': [400],
    'learning_rate': [0.1],
    'max_depth': randint(1, 10),
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'max_features': ['auto', 'sqrt', None]
}

# Create the RandomizedSearchCV object
random_search_gb = RandomizedSearchCV(estimator=GradientBoostingClassifier(), param
random_search_gb
```

```
Out[75]: RandomizedSearchCV
  estimator: GradientBoostingClassifier
    GradientBoostingClassifier
```

Learning hyper-parameters by RandomizedSearchCV

```
In [76]: # # Train the classifier
# random_search_gb.fit(X_train, Y_train)

# best_params = random_search_gb.best_params_
# print("Best Hyperparameters:", best_params)
```

Fitting the Gradient Boosting model by using the learned hyper-parameters

```
In [77]: gb = GradientBoostingClassifier(learning_rate=0.1, max_depth=6, max_features='sqrt')
gb.fit(X_train, Y_train)
y_pred_gb =gb.predict(X_train)
accuracy = accuracy_score(Y_train, y_pred_gb)
print("Train Accuracy with GrBoost Classifier:", accuracy)
```

```
Train Accuracy with GrBoost Classifier: 1.0
```

Saving prediction to CSV

```
In [78]: pred_and_save_to_csv(X_test, gb, 'submission.csv')
```

To take test_input.csv file and give prediction in test_output.csv

```
In [84]: test_input = pd.read_csv('test_input.csv')
X_test_input = test_input

# Performing data-cleaning
X_test_input.drop(['Feature 16', 'Feature 17'], axis=1, inplace=True)
```

```
X_test_input = pd.DataFrame(knn_imputer.transform(X_test_input), columns=X_test_input.columns)
X_test_input.drop(corr_features,axis=1, inplace = True)

# Getting prediction
y_pred =gb.predict(X_test_input)
pred_and_save_to_csv(X_test, gb, 'test_output.csv')
```